

Question 1(a) [3 marks]

Define Machine Learning using suitable example?

Answer:

Machine Learning is a subset of artificial intelligence that enables computers to learn and make decisions from data without being explicitly programmed for every task.

Table: Key Components of Machine Learning

Component	Description
Data	Input information used for training
Algorithm	Mathematical model that learns patterns
Training	Process of teaching the algorithm
Prediction	Output based on learned patterns

Example: Email spam detection system learns from thousands of emails labeled as "spam" or "not spam" to automatically classify new emails.


Mnemonic: "Data Drives Decisions" - Data trains algorithms to make intelligent decisions

Question 1(b) [4 marks]

Explain the process of machine learning with the help of schematic representation

Answer:

The machine learning process involves systematic steps from data collection to model deployment.



Process Steps:

- **Data Collection:** Gathering relevant dataset
- **Preprocessing:** Cleaning and preparing data
- **Training:** Teaching algorithm using training data
- **Validation:** Testing model performance
- **Deployment:** Using model for real predictions

Mnemonic: "Computers Can Truly Think" - Collect, Clean, Train, Test

Question 1(c) [7 marks]

Explain different types of machine learning with suitable application.

Answer:

Machine learning algorithms are categorized based on learning approach and available data.

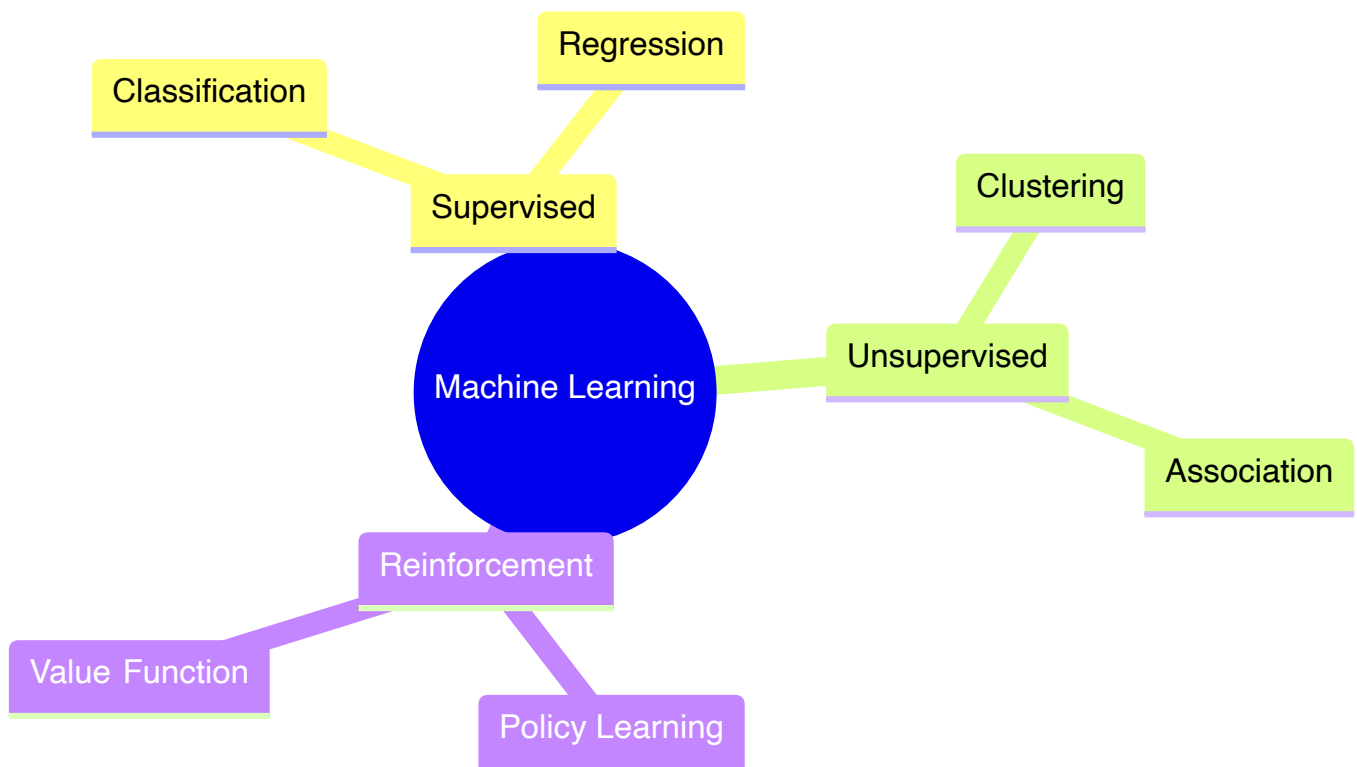
Table: Types of Machine Learning

Type	Learning Method	Data Requirement	Example Application
Supervised	Uses labeled data	Input-output pairs	Email classification
Unsupervised	Finds hidden patterns	Only input data	Customer segmentation
Reinforcement	Learns through rewards	Environment feedback	Game playing AI

Applications:

- **Supervised Learning:** Medical diagnosis, image recognition, fraud detection
- **Unsupervised Learning:** Market research, anomaly detection, recommendation systems
- **Reinforcement Learning:** Autonomous vehicles, robotics, strategic games

Diagram: Learning Types



Mnemonic: "Students Usually Remember" - Supervised, Unsupervised, Reinforcement

Question 1(c) OR [7 marks]

What are various issues with machine learning? List three problems that are not to be solved using machine learning.

Answer:

Table: Machine Learning Issues

Issue Category	Description	Impact
Data Quality	Incomplete, noisy, biased data	Poor model performance
Overfitting	Model memorizes training data	Poor generalization
Computational	High processing requirements	Resource constraints
Interpretability	Black box models	Lack of transparency

Problems NOT suitable for ML:

1. **Simple rule-based tasks** - Basic calculations, simple if-then logic
2. **Ethical decisions** - Moral judgments requiring human values
3. **Creative expression** - Original artistic creation requiring human emotion

Other Issues:

- **Privacy concerns:** Sensitive data handling
- **Bias propagation:** Unfair algorithmic decisions
- **Feature selection:** Choosing relevant input variables

Mnemonic: "Data Drives Quality" - Data quality directly affects model quality

Question 2(a) [3 marks]

Give a summarized view of different types of data in a typical machine learning problem.

Answer:

Table: Data Types in Machine Learning

Data Type	Description	Example
Numerical	Quantitative values	Age: 25, Height: 170cm
Categorical	Discrete categories	Color: Red, Blue, Green
Ordinal	Ordered categories	Rating: Poor, Good, Excellent
Binary	Two possible values	Gender: Male/Female

Characteristics:

- **Structured:** Organized in tables (databases, spreadsheets)
- **Unstructured:** Images, text, audio files
- **Time-series:** Data points over time

Mnemonic: "Numbers Count Better Than Words" - Numerical, Categorical, Binary, Text

Question 2(b) [4 marks]

Calculate variance for both attributes. Determine which attribute is spread out around mean.

Answer:

Given Data:

- Attribute 1: 32, 37, 47, 50, 59
- Attribute 2: 48, 40, 41, 47, 49

Calculations:

Attribute 1:

- Mean = $(32+37+47+50+59)/5 = 225/5 = 45$
- Variance = $[(32-45)^2 + (37-45)^2 + (47-45)^2 + (50-45)^2 + (59-45)^2]/5$
- Variance = $[169 + 64 + 4 + 25 + 196]/5 = 458/5 = 91.6$

Attribute 2:

- Mean = $(48+40+41+47+49)/5 = 225/5 = 45$
- Variance = $[(48-45)^2 + (40-45)^2 + (41-45)^2 + (47-45)^2 + (49-45)^2]/5$
- Variance = $[9 + 25 + 16 + 4 + 16]/5 = 70/5 = 14$

Result: Attribute 1 (variance = 91.6) is more spread out than Attribute 2 (variance = 14).

Mnemonic: "Higher Variance Shows Spread" - Greater variance indicates more dispersion

Question 2(c) [7 marks]

List Factors that lead to data quality issue. How to handle outliers and missing values.

Answer:

Table: Data Quality Issues

Factor	Cause	Solution
Incompleteness	Missing data collection	Imputation techniques
Inconsistency	Different data formats	Standardization
Inaccuracy	Human/sensor errors	Validation rules
Noise	Random variations	Filtering methods

Handling Outliers:

- **Detection:** Statistical methods (Z-score, IQR)

- **Treatment:** Remove, transform, or cap extreme values
- **Visualization:** Box plots, scatter plots

Handling Missing Values:

- **Deletion:** Remove incomplete records
- **Imputation:** Fill with mean, median, or mode
- **Prediction:** Use ML to predict missing values

Code Example:

```
# Handle missing values
df.fillna(df.mean()) # Mean imputation
df.dropna()          # Remove missing rows
```

Mnemonic: "Clean Data Makes Models" - Clean data produces better models

Question 2(a) OR [3 marks]

Give different machine learning activities.

Answer:

Table: Machine Learning Activities

Activity	Purpose	Example
Data Collection	Gather relevant information	Surveys, sensors, databases
Data Preprocessing	Clean and prepare data	Remove noise, handle missing values
Feature Engineering	Create meaningful variables	Extract features from raw data
Model Training	Teach algorithm patterns	Use training dataset
Model Evaluation	Assess performance	Test accuracy, precision, recall
Model Deployment	Put model into production	Web services, mobile apps

Key Activities:

- **Exploratory Data Analysis:** Understanding data patterns
- **Hyperparameter Tuning:** Optimizing model settings
- **Cross-validation:** Robust performance assessment

Mnemonic: "Data Models Perform Excellently" - Data preparation, Model building, Performance evaluation, Execution

Question 2(b) OR [4 marks]

Calculate mean and median of the following numbers: 12,15,18,20,22,24,28,30

Answer:

Given numbers: 12, 15, 18, 20, 22, 24, 28, 30

Mean Calculation:

Mean = $(12+15+18+20+22+24+28+30)/8 = 169/8 = 21.125$

Median Calculation:

- Numbers are already sorted: 12, 15, 18, 20, 22, 24, 28, 30
- Even count (8 numbers)
- Median = $(4\text{th number} + 5\text{th number})/2 = (20 + 22)/2 = 21$

Table: Statistical Summary

Measure	Value	Description
Mean	21.125	Average value
Median	21	Middle value
Count	8	Total numbers

Mnemonic: "Middle Makes Median" - Middle value gives median

Question 2(c) OR [7 marks]

Write a short note on dimensionality reduction and feature subset selection in context with data preprocessing.

Answer:

Dimensionality Reduction removes irrelevant features and reduces computational complexity while preserving important information.

Table: Dimensionality Reduction Techniques

Technique	Method	Use Case
PCA	Principal Component Analysis	Linear reduction
LDA	Linear Discriminant Analysis	Classification tasks
t-SNE	Non-linear embedding	Visualization
Feature Selection	Select important features	Reduce overfitting

Feature Subset Selection Methods:

- **Filter Methods:** Statistical tests, correlation analysis
- **Wrapper Methods:** Forward/backward selection
- **Embedded Methods:** LASSO, Ridge regression

Benefits:

- **Computational Efficiency:** Faster training and prediction
- **Storage Reduction:** Less memory requirements
- **Noise Reduction:** Remove irrelevant features
- **Visualization:** Enable 2D/3D plotting

Code Example:

```
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
reduced_data = pca.fit_transform(data)
```

Mnemonic: "Reduce Features, Improve Performance" - Fewer features often lead to better models

Question 3(a) [3 marks]

Does bias affect the performance of the ML model? Explain briefly.

Answer:

Yes, bias significantly affects ML model performance by creating systematic errors in predictions.

Table: Types of Bias

Bias Type	Description	Impact
Selection Bias	Non-representative data	Poor generalization
Confirmation Bias	Favoring expected results	Skewed conclusions
Algorithmic Bias	Model assumptions	Unfair predictions

Effects on Performance:

- **Underfitting:** High bias leads to oversimplified models
- **Poor Accuracy:** Systematic errors reduce overall performance
- **Unfair Decisions:** Biased models discriminate against groups

Mitigation Strategies:

- Diverse training data
- Cross-validation techniques
- Bias detection algorithms

Mnemonic: "Bias Breaks Better Performance" - Bias reduces model effectiveness

Question 3(b) [4 marks]

Compare cross-validation and bootstrap sampling

Answer:

Table: Cross-validation vs Bootstrap Sampling

Aspect	Cross-validation	Bootstrap Sampling
Method	Split data into folds	Sample with replacement
Data Usage	Uses all data	Creates multiple samples
Purpose	Model evaluation	Estimate uncertainty
Overlap	No overlap between sets	Allows duplicate samples

Cross-validation:

- Divides data into k equal parts
- Trains on $k-1$ parts, tests on 1 part
- Repeats k times for robust evaluation

Bootstrap Sampling:

- Creates random samples with replacement
- Generates multiple datasets of same size
- Estimates confidence intervals

Applications:

- **Cross-validation:** Model selection, hyperparameter tuning
- **Bootstrap:** Statistical inference, confidence estimation

Mnemonic: "Cross Checks, Bootstrap Builds" - Cross-validation checks performance, Bootstrap builds confidence

Question 3(c) [7 marks]

Confusion Matrix Calculation and Metrics

Answer:

Given Information:

- True Positive (TP): 83 (predicted buy, actually bought)
- False Positive (FP): 7 (predicted buy, didn't buy)
- False Negative (FN): 5 (predicted no buy, actually bought)
- True Negative (TN): 5 (predicted no buy, didn't buy)

Confusion Matrix:

	Predicted Buy	Predicted No Buy
Actually Buy	83 (TP)	5 (FN)
Actually No Buy	7 (FP)	5 (TN)

Calculations:**a) Error Rate:**

$$\text{Error Rate} = (\text{FP} + \text{FN}) / \text{Total} = (7 + 5) / 100 = 0.12 = 12\%$$

b) Precision:

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP}) = 83 / (83 + 7) = 83/90 = 0.922 = 92.2\%$$

c) Recall:

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN}) = 83 / (83 + 5) = 83/88 = 0.943 = 94.3\%$$

d) F-measure:

$$\text{F-measure} = 2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$$

$$\text{F-measure} = 2 \times (0.922 \times 0.943) / (0.922 + 0.943) = 0.932 = 93.2\%$$

Table: Performance Metrics

Metric	Value	Interpretation
Error Rate	12%	Model makes 12% wrong predictions
Precision	92.2%	92.2% of predicted buyers actually buy
Recall	94.3%	Model identifies 94.3% of actual buyers
F-measure	93.2%	Balanced performance measure

Mnemonic: "Perfect Recall Finds Everyone" - Precision measures accuracy, Recall finds all positives

Question 3(a) OR [3 marks]

Define in brief: a) Target function b) Cost function c) Loss Function

Answer:

Table: Function Definitions

Function	Definition	Purpose
Target Function	Ideal mapping from input to output	What we want to learn
Cost Function	Measures overall model error	Evaluate total performance
Loss Function	Measures error for single prediction	Individual prediction error

Detailed Explanation:

- **Target Function:** $f(x) = y$, the true relationship we want to approximate
- **Cost Function:** Average of all loss functions, $J = (1/n)\sum \text{loss}(y_i, \hat{y}_i)$
- **Loss Function:** Error for one sample, e.g., $(y_i - \hat{y}_i)^2$

Relationship: Cost function is typically the average of loss functions across all training examples.

Mnemonic: "Target Costs Less" - Target function is ideal, Cost function measures overall error, Loss function measures individual error

Question 3(b) OR [4 marks]

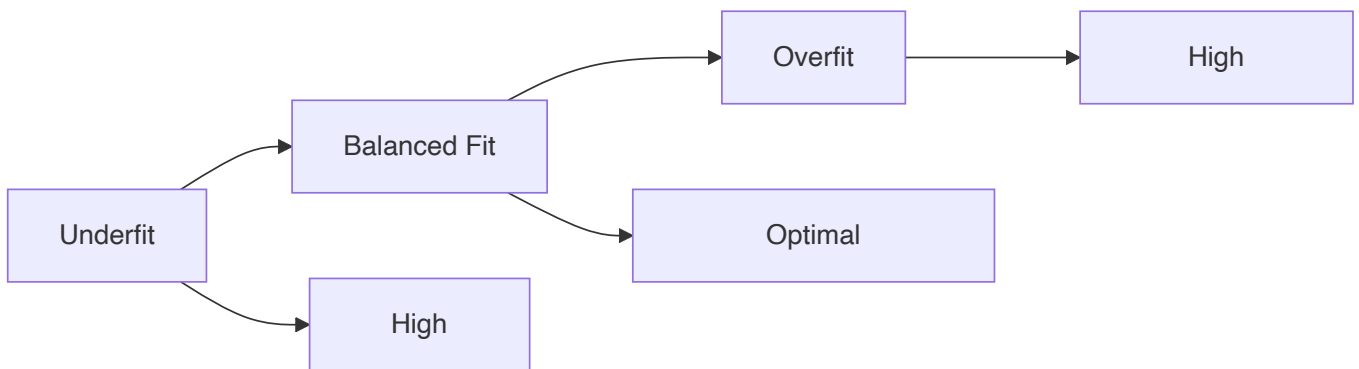
Explain balanced fit, underfit and overfit

Answer:

Table: Model Fitting Types

Fit Type	Training Error	Validation Error	Characteristics
Underfit	High	High	Too simple model
Balanced Fit	Low	Low	Optimal complexity
Overfit	Very Low	High	Too complex model

Visualization:



Characteristics:

- **Underfit:** Model too simple, cannot capture patterns
- **Balanced Fit:** Right complexity, generalizes well
- **Overfit:** Model too complex, memorizes training data

Solutions:

- **Underfit:** Increase model complexity, add features
- **Overfit:** Regularization, cross-validation, more data

Mnemonic: "Balance Brings Best Results" - Balanced models perform best on new data

Question 4(a) [3 marks]

Give classification learning steps.

Answer:

Table: Classification Learning Steps

Step	Description	Purpose
Data Collection	Gather labeled examples	Provide training material
Preprocessing	Clean and prepare data	Improve data quality
Feature Selection	Choose relevant attributes	Reduce complexity
Model Training	Learn from training data	Build classifier
Evaluation	Test model performance	Assess accuracy
Deployment	Use for new predictions	Practical application

Detailed Process:

- 1. **Prepare dataset** with input features and class labels
- 2. **Split data** into training and testing sets
- 3. **Train classifier** using training data
- 4. **Validate model** using test data
- 5. **Fine-tune parameters** for optimal performance

Mnemonic: "Data Preparation Facilitates Model Excellence" - Data prep, Feature selection, Model training, Evaluation

Question 4(b) [4 marks]

Linear Relationship Calculation

Answer:

Given Data:

Hours (X)	Exam Score (Y)
2	85
3	80
4	75
5	70
6	60

Linear Regression Calculation:**Step 1: Calculate means**

- $\bar{X} = (2+3+4+5+6)/5 = 4$
- $\bar{Y} = (85+80+75+70+60)/5 = 74$

Step 2: Calculate slope (b)

- Numerator = $\sum(X-\bar{X})(Y-\bar{Y}) = (2-4)(85-74) + (3-4)(80-74) + (4-4)(75-74) + (5-4)(70-74) + (6-4)(60-74)$
- $= (-2)(11) + (-1)(6) + (0)(1) + (1)(-4) + (2)(-14) = -22 - 6 + 0 - 4 - 28 = -60$
- Denominator = $\sum(X-\bar{X})^2 = (-2)^2 + (-1)^2 + (0)^2 + (1)^2 + (2)^2 = 4 + 1 + 0 + 1 + 4 = 10$
- $b = -60/10 = -6$

Step 3: Calculate intercept (a)

- $a = \bar{Y} - b \times \bar{X} = 74 - (-6) \times 4 = 74 + 24 = 98$

Linear Equation: $Y = 98 - 6X$

Interpretation: For every additional hour of smartphone use, exam score decreases by 6 points.

Mnemonic: "More Phone, Less Score" - Negative correlation between phone use and grades

Question 4(c) [7 marks]

Explain classification steps in detail**Answer:**

Classification is a supervised learning process that assigns input data to predefined categories or classes.

Detailed Classification Steps:**1. Problem Definition**

- Define classes and objectives
- Identify input features and target variable
- Determine success criteria

2. Data Collection and Preparation



3. Feature Engineering

- **Feature Selection:** Choose relevant attributes
- **Feature Extraction:** Create new meaningful features
- **Normalization:** Scale features to similar ranges

4. Model Selection and Training

Table: Common Classification Algorithms

Algorithm	Best For	Advantages
Decision Tree	Interpretable rules	Easy to understand
SVM	High-dimensional data	Good generalization
Neural Networks	Complex patterns	High accuracy
Naive Bayes	Text classification	Fast training

5. Model Evaluation

- **Confusion Matrix:** Detailed performance analysis
- **Cross-validation:** Robust performance estimation
- **Metrics:** Accuracy, Precision, Recall, F1-score

6. Hyperparameter Tuning

- Grid search for optimal parameters
- Validation set for parameter selection

7. Final Evaluation and Deployment

- Test on unseen data
- Deploy model for production use
- Monitor performance over time

Mnemonic: "Proper Data Modeling Evaluates Performance Thoroughly" - Problem definition, Data prep, Modeling, Evaluation, Performance testing, Tuning

Question 4(a) OR [3 marks]

Does the choice of the k value influence the performance of the KNN algorithm? Explain briefly

Answer:

Yes, the k value significantly influences KNN algorithm performance by affecting the decision boundary and model complexity.

Table: K Value Impact

K Value	Effect	Performance
Small K (k=1)	Sensitive to noise	High variance, low bias
Medium K	Balanced decisions	Optimal performance
Large K	Smooth boundaries	Low variance, high bias

Impact Analysis:

- **k=1:** May overfit to training data, sensitive to outliers
- **Optimal k:** Usually odd number, balances bias-variance tradeoff
- **Large k:** May underfit, loses local patterns

Selection Strategy:

- Use cross-validation to find optimal k
- Try $k = \sqrt{n}$ as starting point
- Consider computational cost vs accuracy

Mnemonic: "Small K Varies, Large K Smooths" - Small k creates variance, large k creates smooth boundaries

Question 4(b) OR [4 marks]

Define Support Vectors in the SVM model.

Answer:

Support Vectors are the critical data points that lie closest to the decision boundary (hyperplane) in Support Vector Machine algorithm.

Table: Support Vector Characteristics

Aspect	Description	Importance
Location	Closest points to hyperplane	Define decision boundary
Distance	Equal distance from boundary	Maximize margin
Role	Support the hyperplane	Determine optimal separation
Sensitivity	Removing them changes model	Critical for model structure

Key Properties:

- **Margin Definition:** Support vectors determine the maximum margin between classes
- **Model Dependency:** Only support vectors affect the final model

- **Boundary Formation:** Create the optimal separating hyperplane

Diagram:

```

Class A |   | Class B
o       |   | x
  o     |   | x
o  O    |   | X  x
  o     |   | x
o       |   | x

Support Vectors: O and X
Hyperplane: ---|---

```

Mathematical Significance: Support vectors satisfy the constraint $y_i(w \cdot x_i + b) = 1$, where they lie exactly on the margin boundary.

Mnemonic: "Support Vectors Support Decisions" - These vectors support the decision boundary

Question 4(c) OR [7 marks]

Explain logistic regression in detail.

Answer:

Logistic Regression is a statistical method used for binary classification that models the probability of class membership using the logistic function.

Mathematical Foundation:

Sigmoid Function:

$$\sigma(z) = 1 / (1 + e^{(-z)})$$

$$\text{where } z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Table: Linear vs Logistic Regression

Aspect	Linear Regression	Logistic Regression
Output	Continuous values	Probabilities (0-1)
Function	Linear	Sigmoid (S-curve)
Purpose	Prediction	Classification
Error Function	Mean Squared Error	Log-likelihood

Key Components:

1. Logistic Function Properties:

- **S-shaped curve:** Smooth transition between 0 and 1
- **Asymptotes:** Approaches 0 and 1 but never reaches them

- **Monotonic:** Always increasing function

2. Model Training:

- **Maximum Likelihood Estimation:** Find parameters that maximize probability of observed data
- **Gradient Descent:** Iterative optimization algorithm
- **Cost Function:** Log-loss or cross-entropy

3. Decision Making:

- **Threshold:** Typically 0.5 for binary classification
- **Probability Output:** $P(y=1 | x)$ gives class probability
- **Decision Rule:** Classify as positive if $P(y=1 | x) > 0.5$

Advantages:

- **Probabilistic Output:** Provides confidence in predictions
- **No Assumptions:** About distribution of independent variables
- **Less Overfitting:** Compared to complex models
- **Fast Training:** Efficient computation

Applications:

- Medical diagnosis
- Marketing response prediction
- Credit approval decisions
- Email spam detection

Code Example:

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
predictions = model.predict(X_test)
probabilities = model.predict_proba(X_test)
```

Mnemonic: "Sigmoid Squashes Infinite Input" - Sigmoid function converts any real number to probability

Question 5(a) [3 marks]

Write a short note on Matplotlib python library.

Answer:

Matplotlib is a comprehensive Python library for creating static, animated, and interactive visualizations in data science and machine learning.

Table: Matplotlib Key Features

Feature	Purpose	Example
Pyplot	MATLAB-like plotting interface	Line plots, scatter plots
Object-oriented	Advanced customization	Figure and axes objects
Multiple formats	Save in various formats	PNG, PDF, SVG, EPS
Subplots	Multiple plots in one figure	Grid arrangements

Common Plot Types:

- **Line Plot:** Trends over time
- **Scatter Plot:** Relationship between variables
- **Histogram:** Data distribution
- **Bar Chart:** Categorical comparisons
- **Box Plot:** Statistical summaries

Basic Usage:

```
import matplotlib.pyplot as plt
plt.plot(x, y)
plt.xlabel('X Label')
plt.ylabel('Y Label')
plt.title('Plot Title')
plt.show()
```

Applications: Data exploration, model performance visualization, presentation graphics

Mnemonic: "Matplotlib Makes Pretty Plots" - Essential tool for data visualization

Question 5(b) [4 marks]

K-means clustering for two-dimensional data

Answer:

Given Points:

{(2,3),(3,3),(4,3),(5,3),(6,3),(7,3),(8,3),(25,20),(26,20),(27,20),(28,20),(29,20),(30,20)}

K-means Algorithm Steps:**Step 1: Initialize centroids**

- Cluster 1: (4, 3) - chosen from left group
- Cluster 2: (27, 20) - chosen from right group

Step 2: Assign points to nearest centroid**Table: Point Assignments**

Point	Distance to C1	Distance to C2	Assigned Cluster
(2,3)	2.0	25.8	Cluster 1
(3,3)	1.0	24.8	Cluster 1
(4,3)	0.0	23.8	Cluster 1
(5,3)	1.0	22.8	Cluster 1
(6,3)	2.0	21.8	Cluster 1
(7,3)	3.0	20.8	Cluster 1
(8,3)	4.0	19.8	Cluster 1
(25,20)	23.8	2.0	Cluster 2
(26,20)	24.8	1.0	Cluster 2
(27,20)	25.8	0.0	Cluster 2
(28,20)	26.8	1.0	Cluster 2
(29,20)	27.8	2.0	Cluster 2
(30,20)	28.8	3.0	Cluster 2

Step 3: Update centroids

- New C1 = $((2+3+4+5+6+7+8)/7, (3+3+3+3+3+3+3)/7) = (5, 3)$
- New C2 = $((25+26+27+28+29+30)/6, (20+20+20+20+20+20)/6) = (27.5, 20)$

Final Clusters:

- **Cluster 1:** $\{(2,3),(3,3),(4,3),(5,3),(6,3),(7,3),(8,3)\}$
- **Cluster 2:** $\{(25,20),(26,20),(27,20),(28,20),(29,20),(30,20)\}$

Mnemonic: "Centroids Attract Nearest Neighbors" - Points join closest centroid

Question 5(c) [7 marks]

Give functions and its use of Scikit-learn for: a. Data Preprocessing b. Model Selection c. Model Evaluation and Metrics

Answer:

Scikit-learn provides comprehensive tools for machine learning workflow from data preprocessing to model evaluation.

a) Data Preprocessing Functions:**Table: Preprocessing Functions**

Function	Purpose	Example Usage
<code>StandardScaler()</code>	Normalize features	Remove mean, unit variance
<code>MinMaxScaler()</code>	Scale to range [0,1]	Feature scaling
<code>LabelEncoder()</code>	Encode categorical labels	Convert text to numbers
<code>OneHotEncoder()</code>	Create dummy variables	Handle categorical features
<code>train_test_split()</code>	Split dataset	Training/testing division

Code Example:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

b) Model Selection Functions:**Table: Model Selection Tools**

Function	Purpose	Application
<code>GridSearchCV()</code>	Hyperparameter tuning	Find optimal parameters
<code>RandomizedSearchCV()</code>	Random parameter search	Faster parameter optimization
<code>cross_val_score()</code>	Cross-validation	Model performance evaluation
<code>StratifiedKFold()</code>	Stratified sampling	Balanced cross-validation
<code>Pipeline()</code>	Combine preprocessing and modeling	Streamlined workflow

Code Example:

```
from sklearn.model_selection import GridSearchCV
param_grid = {'C': [0.1, 1, 10]}
grid_search = GridSearchCV(SVM(), param_grid, cv=5)
grid_search.fit(X_train, y_train)
```

c) Model Evaluation and Metrics Functions:**Table: Evaluation Metrics**

Function	Purpose	Use Case
<code>accuracy_score()</code>	Overall accuracy	General classification
<code>precision_score()</code>	Positive prediction accuracy	Minimize false positives
<code>recall_score()</code>	True positive rate	Minimize false negatives
<code>f1_score()</code>	Harmonic mean of precision/recall	Balanced metric
<code>confusion_matrix()</code>	Detailed error analysis	Understanding mistakes
<code>classification_report()</code>	Comprehensive metrics	Complete evaluation
<code>roc_auc_score()</code>	Area under ROC curve	Binary classification

Code Example:

```
from sklearn.metrics import classification_report
print(classification_report(y_true, y_pred))
```

Workflow Integration:

- **Preprocessing:** Clean and prepare data
- **Model Selection:** Choose and tune algorithms
- **Evaluation:** Assess performance comprehensively

Mnemonic: "Preprocess, Select, Evaluate" - Complete ML workflow in Scikit-learn

Question 5(a) OR [3 marks]

List out the major features of Numpy.

Answer:

NumPy (Numerical Python) is the fundamental package for scientific computing in Python, providing powerful array operations and mathematical functions.

Table: Major NumPy Features

Feature	Description	Benefit
N-dimensional Arrays	Efficient array objects	Fast mathematical operations
Broadcasting	Operations on different sized arrays	Flexible computations
Linear Algebra	Matrix operations, decompositions	Scientific computing
Random Numbers	Random sampling and distributions	Statistical simulations
Integration	Works with C/C++/Fortran	High performance

Key Capabilities:

- **Mathematical Functions:** Trigonometric, logarithmic, exponential
- **Array Manipulation:** Reshaping, splitting, joining arrays
- **Indexing:** Advanced slicing and boolean indexing
- **Memory Efficiency:** Optimized data storage

Applications: Data analysis, machine learning, image processing, scientific research

Mnemonic: "Numbers Need Numpy's Power" - Essential for numerical computations

Question 5(b) OR [4 marks]

K-means clustering for one-dimensional data

Answer:

Given Dataset: {1,2,4,5,7,8,10,11,12,14,15,17}

K-means Algorithm for 3 clusters:

Step 1: Initialize centroids

- C1 = 3 (around early values)
- C2 = 9 (around middle values)
- C3 = 15 (around later values)

Step 2: Assign points to nearest centroid

Table: Point Assignments (Iteration 1)

Point	Distance to C1	Distance to C2	Distance to C3	Assigned Cluster
1	2	8	14	Cluster 1
2	1	7	13	Cluster 1
4	1	5	11	Cluster 1
5	2	4	10	Cluster 1
7	4	2	8	Cluster 2
8	5	1	7	Cluster 2
10	7	1	5	Cluster 2
11	8	2	4	Cluster 2
12	9	3	3	Cluster 2
14	11	5	1	Cluster 3
15	12	6	0	Cluster 3
17	14	8	2	Cluster 3

Step 3: Update centroids

- New C1 = $(1+2+4+5)/4 = 3$
- New C2 = $(7+8+10+11+12)/5 = 9.6$
- New C3 = $(14+15+17)/3 = 15.33$

Final Clusters:

- **Cluster 1:** {1, 2, 4, 5}
- **Cluster 2:** {7, 8, 10, 11, 12}
- **Cluster 3:** {14, 15, 17}

Mnemonic: "Groups Gather by Distance" - Similar points form natural clusters

Question 5(c) OR [7 marks]

Give function and its use of Pandas library for: a. Data Preprocessing b. Data Inspection c. Data Cleaning and Transformation

Answer:

Pandas is a powerful Python library for data manipulation and analysis, providing high-level data structures and operations.

a) Data Preprocessing Functions:

Table: Preprocessing Functions

Function	Purpose	Example
<code>read_csv()</code>	Load CSV files	<code>pd.read_csv('data.csv')</code>
<code>head()</code>	View first n rows	<code>df.head(10)</code>
<code>tail()</code>	View last n rows	<code>df.tail(5)</code>
<code>sample()</code>	Random sampling	<code>df.sample(100)</code>
<code>set_index()</code>	Set column as index	<code>df.set_index('id')</code>

b) Data Inspection Functions:

Table: Inspection Functions

Function	Purpose	Information Provided
<code>info()</code>	Dataset overview	Data types, memory usage
<code>describe()</code>	Statistical summary	Mean, std, min, max
<code>shape</code>	Dataset dimensions	(rows, columns)
<code>dtypes</code>	Data types	Column data types
<code>isnull()</code>	Missing values	Boolean mask for nulls
<code>value_counts()</code>	Count unique values	Frequency distribution
<code>corr()</code>	Correlation matrix	Feature relationships

Code Example:

```
# Data inspection
print(df.info())
print(df.describe())
print(df.isnull().sum())
```

c) Data Cleaning and Transformation Functions:

Table: Cleaning Functions

Function	Purpose	Usage
<code>dropna()</code>	Remove missing values	<code>df.dropna()</code>
<code>fillna()</code>	Fill missing values	<code>df.fillna(0)</code>
<code>drop_duplicates()</code>	Remove duplicate rows	<code>df.drop_duplicates()</code>
<code>replace()</code>	Replace values	<code>df.replace('old', 'new')</code>
<code>astype()</code>	Change data types	<code>df['col'].astype('int')</code>
<code>apply()</code>	Apply function to data	<code>df.apply(lambda x: x*2)</code>
<code>groupby()</code>	Group data	<code>df.groupby('category')</code>
<code>merge()</code>	Join datasets	<code>pd.merge(df1, df2)</code>
<code>pivot()</code>	Reshape data	<code>df.pivot(columns='col')</code>

Advanced Operations:

- **String Operations:** `str.contains()`, `str.replace()`
- **Date Operations:** `to_datetime()`, `dt.year`
- **Categorical Data:** `pd.Categorical()`

Workflow Example:

```
# Complete preprocessing pipeline
df = pd.read_csv('data.csv')
df = df.dropna()
df['category'] = df['category'].astype('category')
df_grouped = df.groupby('type').mean()
```

Benefits:

- **Intuitive Syntax:** Easy to learn and use
- **Performance:** Optimized for large datasets
- **Integration:** Works well with NumPy, Matplotlib
- **Flexibility:** Handles various data formats

Mnemonic: "Pandas Processes Data Perfectly" - Comprehensive data manipulation tool