

Question 1(a) [3 marks]

Define following terms: 1) Fuzzy Logic. 2) Expert System.

Answer:

Term	Definition
Fuzzy Logic	Computing approach that deals with approximate rather than fixed and exact reasoning, allowing degrees of truth between 0 and 1
Expert System	AI program that mimics human expert decision-making by using knowledge base and inference engine to solve specific problems

- **Key Features:** Both handle uncertainty and incomplete information
- **Applications:** Medical diagnosis, industrial control systems

Mnemonic: "Fuzzy Experts handle uncertain decisions"

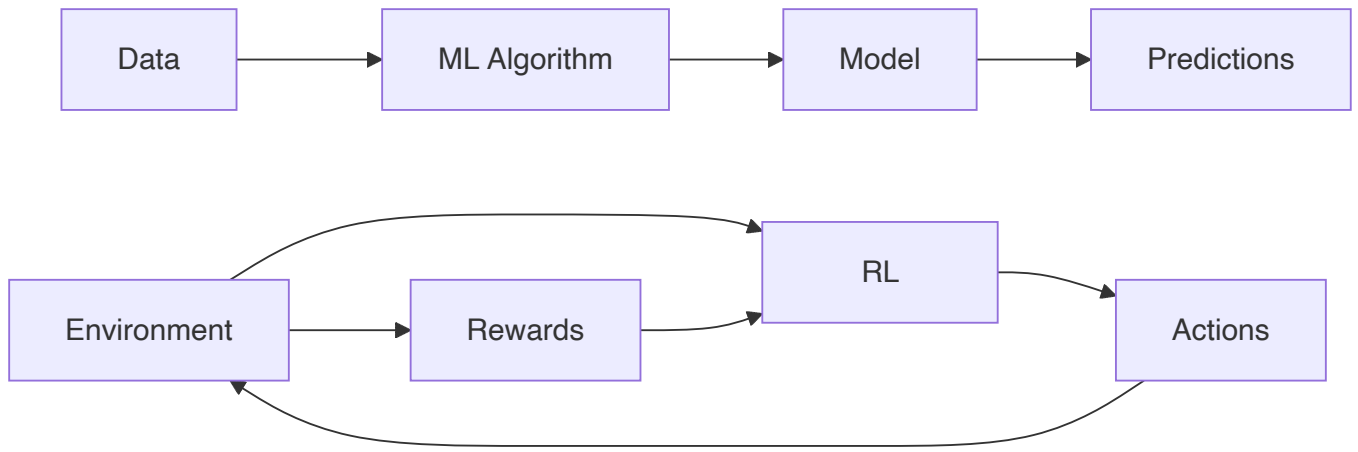
Question 1(b) [4 marks]

Define following terms: 1) Machine Learning. 2) Reinforcement Learning.

Answer:

Term	Definition	Key Characteristic
Machine Learning	Algorithm that automatically improves performance through experience without explicit programming	Learning from data patterns
Reinforcement Learning	Agent learns optimal actions through trial-and-error interactions with environment using rewards/penalties	Learning through feedback

Diagram:



Mnemonic: "ML learns from Data, RL learns from Rewards"

Question 1(c) [7 marks]

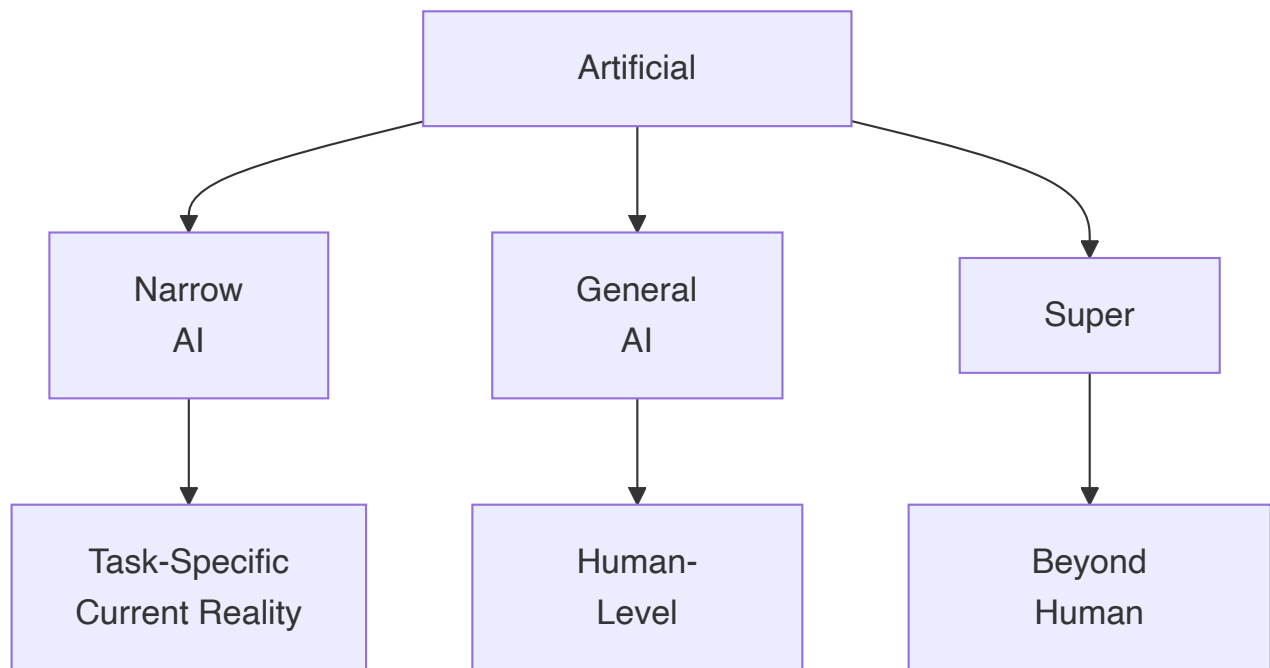
Explain types of Artificial Intelligence in detail with suitable diagram.

Answer:

Table: Types of AI

Type	Description	Capability	Examples
Narrow AI	Designed for specific tasks	Limited domain expertise	Siri, Chess programs
General AI	Human-level intelligence across domains	Multi-domain reasoning	Currently theoretical
Super AI	Exceeds human intelligence	Beyond human capabilities	Future concept

Diagram:



- **Current Status:** We are in Narrow AI era
- **Development Path:** Narrow → General → Super AI
- **Timeline:** General AI expected in 20-30 years

Mnemonic: "Narrow Now, General Goal, Super Soon"

Question 1(c) OR [7 marks]

Explain various aspects related to ethics while designing an AI system. Also, explain limitations of AI system in detail.

Answer:

AI Ethics Table:

Ethical Aspect	Description	Implementation
Fairness	Avoid bias and discrimination	Diverse training data
Transparency	Explainable AI decisions	Clear algorithms
Privacy	Protect user data	Data encryption
Accountability	Responsibility for AI actions	Human oversight

AI Limitations:

- **Data Dependency:** Requires large, quality datasets
- **Lack of Common Sense:** Cannot understand context like humans
- **Brittleness:** Fails in unexpected situations

- **Black Box Problem:** Difficult to explain decisions

Mnemonic: "Fair, Transparent, Private, Accountable AI has Data, Common sense, Brittleness, Black box issues"

Question 2(a) [3 marks]

List characteristics of reinforcement learning.

Answer:

Characteristic	Description
Trial-and-Error	Agent learns through experimentation
Reward-Based	Feedback through rewards/penalties
Sequential Decision Making	Actions affect future states
Exploration vs Exploitation	Balance between trying new actions and using known good actions

Mnemonic: "Trial Rewards Sequential Exploration"

Question 2(b) [4 marks]

Explain positive reinforcement and negative reinforcement.

Answer:

Comparison Table:

Type	Definition	Effect	Example
Positive Reinforcement	Adding pleasant stimulus to increase behavior	Increases desired behavior	Giving treat for good performance
Negative Reinforcement	Removing unpleasant stimulus to increase behavior	Increases desired behavior	Stopping alarm when task completed

Key Difference: Both increase behavior, but positive adds reward while negative removes punishment.

Mnemonic: "Positive Adds pleasure, Negative Removes pain"

Question 2(c) [7 marks]

Explain Supervised learning in detail.

Answer:

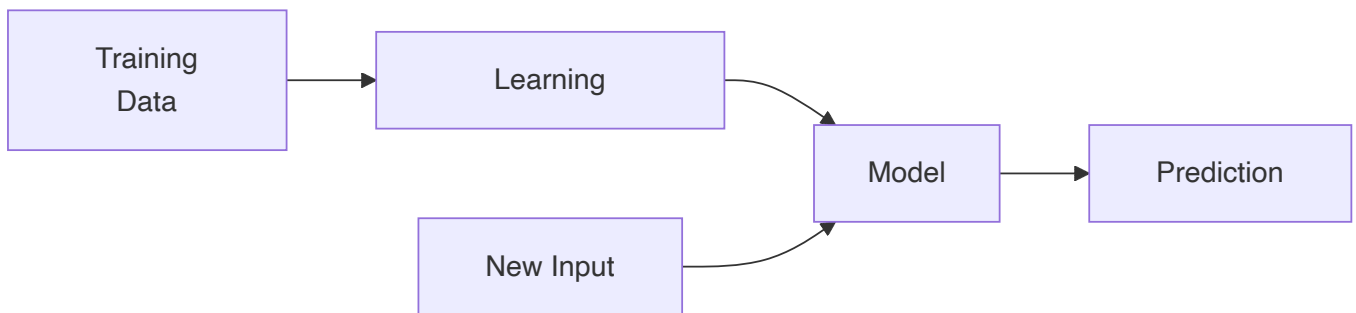
Definition: Learning algorithm that learns from labeled training data to make predictions on new data.

Process Table:

Step	Description	Example
Training	Algorithm learns from input-output pairs	Email → Spam/Not Spam
Validation	Test model on unseen data	Check accuracy
Prediction	Make outputs for new inputs	Classify new emails

Types:

- **Classification:** Predicts categories (spam detection)
- **Regression:** Predicts continuous values (house prices)

Diagram:

Mnemonic: "Supervised = Teacher provides correct answers"

Question 2(a) OR [3 marks]

List key components involved in human learning.

Answer:

Component	Function
Observation	Gathering information from environment
Memory	Storing and retrieving experiences
Practice	Repetition to improve skills
Feedback	Information about performance

Mnemonic: "Observe, Memorize, Practice, Feedback"

Question 2(b) OR [4 marks]

Explain about well-posed learning problem in detail.

Answer:

Definition: A learning problem with clearly defined task, performance measure, and experience.

Components Table:

Component	Description	Example
Task (T)	What the system should learn to do	Play chess
Performance (P)	How to measure success	Win percentage
Experience (E)	Training data or practice	Previous games

Formula: Learning = Improving P on T through E

Criteria: Problem must be measurable, achievable, and have available data.

Mnemonic: "Task Performance Experience = TPE for learning"

Question 2(c) OR [7 marks]

Explain Unsupervised learning in detail.

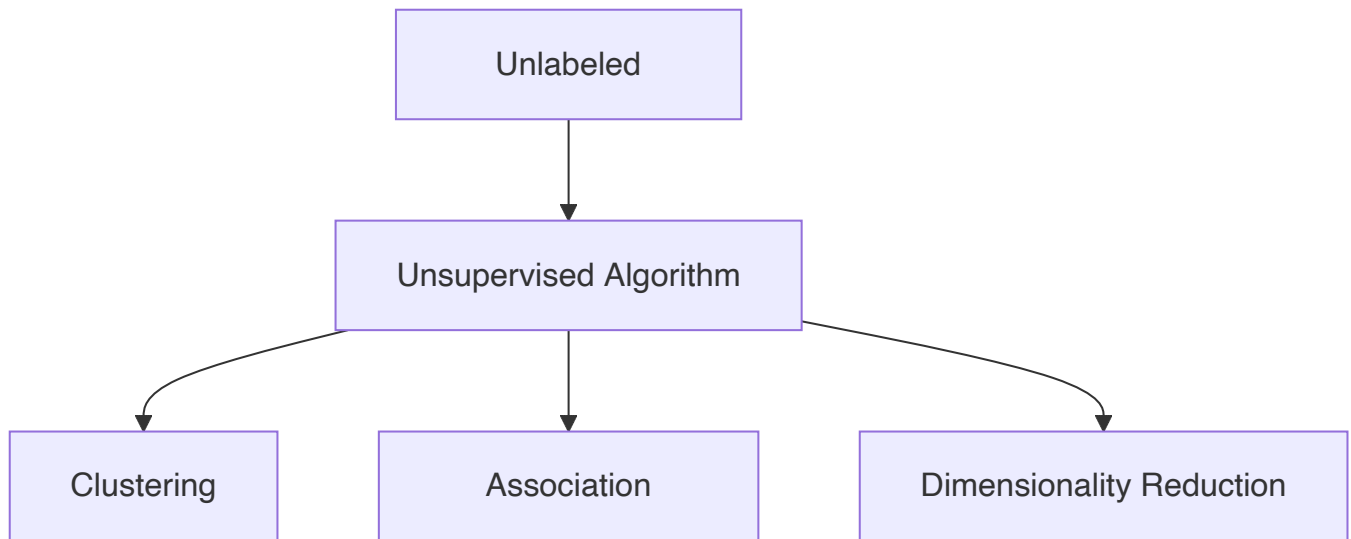
Answer:

Definition: Learning patterns from data without labeled examples or target outputs.

Types Table:

Type	Purpose	Algorithm	Example
Clustering	Group similar data	K-means	Customer segmentation
Association	Find relationships	Apriori	Market basket analysis
Dimensionality Reduction	Reduce features	PCA	Data visualization

Diagram:



- **No Teacher:** Algorithm finds hidden patterns independently
- **Exploratory:** Discovers unknown structures in data

Mnemonic: "Unsupervised = No teacher, find patterns yourself"

Question 3(a) [3 marks]


Explain SIGMOID function. Also, draw its graph and provide an example of SIGMOID function.

Answer:

Definition: Activation function that maps any real number to value between 0 and 1.

Formula: $\sigma(x) = 1/(1 + e^{(-x)})$

Graph (ASCII):



Example: For $x = 0$, $\sigma(0) = 1/(1 + e^0) = 1/2 = 0.5$

Properties: S-shaped curve, smooth gradient, used in binary classification

Mnemonic: "Sigmoid Squashes values between 0 and 1"

Question 3(b) [4 marks]

Define following term: 1) Activation function. 2) Artificial neural network.

Answer:

Term	Definition	Purpose
Activation Function	Mathematical function that determines neuron output based on weighted inputs	Introduces non-linearity to neural networks
Artificial Neural Network	Computing system inspired by biological neural networks with interconnected nodes	Pattern recognition and machine learning

Key Features:

- **Non-linear processing** enables complex pattern learning
- **Layered architecture** processes information hierarchically

Mnemonic: "Activation Artificially mimics brain neurons"

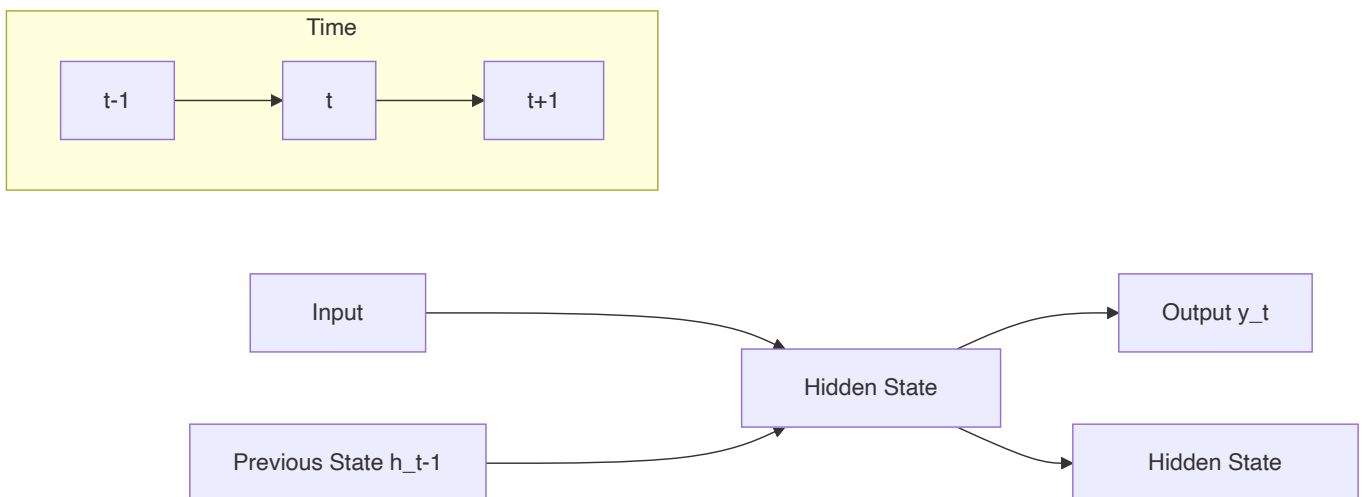
Question 3(c) [7 marks]

Draw and explain architecture of Recurrent network in detail.

Answer:

Definition: Neural network with connections that create loops, allowing information persistence.

Architecture Diagram:



Components Table:

Component	Function	Formula
Hidden State	Memory of previous inputs	$h_t = f(W_h * h_{t-1} + W_x * x_t)$
Input Layer	Current time step input	x_t
Output Layer	Prediction at time t	$y_t = W_y * h_t$

Applications: Speech recognition, language translation, time series prediction

Advantage: Handles sequential data with memory of past information

Mnemonic: "Recurrent = Remembers previous states"

Question 3(a) OR [3 marks]


Explain TANH function. Also, draw its graph and provide an example of TANH function.

Answer:

Definition: Hyperbolic tangent activation function that maps values between -1 and 1.

Formula: $\tanh(x) = (e^x - e^{-x}) / (e^x + e^{-x})$

Graph (ASCII):



Example: For $x = 0$, $\tanh(0) = (1-1)/(1+1) = 0$

Properties: Zero-centered, S-shaped, stronger gradients than sigmoid

Mnemonic: "TANH = Two-sided sigmoid (-1 to +1)"

Question 3(b) OR [4 marks]

Define following term: 1) Biological neural network. 2) Loss calculation.

Answer:

Term	Definition	Key Aspects
Biological Neural Network	Network of interconnected neurons in living organisms that process information	Dendrites, cell body, axon, synapses
Loss calculation	Mathematical measure of difference between predicted and actual outputs	Guides learning through backpropagation

Biological Structure: Neurons → Synapses → Neural Networks → Brain

Loss Types: Mean Squared Error, Cross-entropy, Absolute Error

Mnemonic: "Biology inspires AI, Loss measures learning progress"


Question 3(c) OR [7 marks]

Draw and explain architecture of multi-layer feed-forward network in detail.

Answer:

Definition: Neural network with multiple layers where information flows forward from input to output.

Architecture Diagram:



Layer Functions Table:

Layer	Function	Processing
Input	Receives data	No processing, just distribution
Hidden	Feature extraction	Weighted sum + activation function
Output	Final prediction	Classification or regression output

Information Flow: Input → Hidden Layer(s) → Output (unidirectional)

Learning: Backpropagation adjusts weights based on error

Mnemonic: "Multi-layer = Multiple hidden layers for complex learning"

Question 4(a) [3 marks]

List advantages of NLP in detail.

Answer:

Advantage	Description
Automation	Automates text processing tasks that require human effort
Language Understanding	Processes multiple languages and dialects effectively
24/7 Availability	Works continuously without human intervention
Scalability	Handles large volumes of text data efficiently

Applications: Chatbots, translation, sentiment analysis, document processing

Mnemonic: "NLP = Automates Language Understanding 24/7 at Scale"

Question 4(b) [4 marks]

Explain Natural Language Generation in detail.

Answer:

Definition: AI process that converts structured data into natural human language text.

Process Table:

Step	Description	Function
Content Planning	Decide what information to include	Data selection
Sentence Planning	Structure sentences and paragraphs	Text organization
Surface Realization	Generate actual text with grammar	Final output

Applications: Report generation, chatbots, automated journalism, personalized content

Example: Converting sales data → "Sales increased 15% this quarter due to strong performance in electronics."

Mnemonic: "NLG = Numbers to Narrative"

Question 4(c) [7 marks]

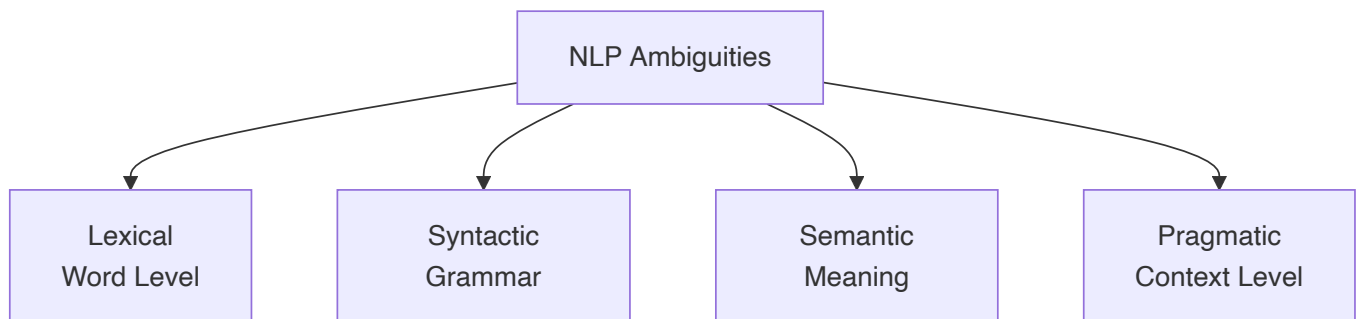
Explain types of ambiguities in NLP. Also, provide examples of each ambiguity.

Answer:

Ambiguity Types Table:

Type	Description	Example	Resolution
Lexical	Word has multiple meanings	"Bank" (river/financial)	Context analysis
Syntactic	Sentence structure unclear	"I saw man with telescope"	Parse trees
Semantic	Meaning unclear	"Colorless green ideas"	Semantic rules
Pragmatic	Context-dependent meaning	"Can you pass salt?" (request/question)	Situational context

Diagram:



Resolution Strategies: Context analysis, statistical models, knowledge bases

Mnemonic: "Lexical Syntactic Semantic Pragmatic = LSSP ambiguities"

Question 4(a) OR [3 marks]

List disadvantages of NLP in detail.

Answer:

Disadvantage	Description
Context Limitations	Struggles with sarcasm, humor, cultural references
Language Complexity	Difficulty with idioms, slang, regional dialects
Data Requirements	Needs large amounts of training data
Computational Cost	Requires significant processing power and memory

Challenges: Ambiguity, multilingual support, real-time processing

Mnemonic: "NLP Challenges = Context, Language, Data, Computation"

Question 4(b) OR [4 marks]

Explain Natural Language Understanding in detail.

Answer:

Definition: AI capability to comprehend and interpret human language meaning and intent.

Components Table:

Component	Function	Example
Tokenization	Break text into words/phrases	"Hello world" → ["Hello", "world"]
Parsing	Analyze grammatical structure	Identify subject, verb, object
Semantic Analysis	Extract meaning	Understand relationships between words
Intent Recognition	Identify user purpose	"Book flight" → travel booking intent

Process Flow: Text Input → Tokenization → Parsing → Semantic Analysis → Intent Understanding

Applications: Virtual assistants, chatbots, voice commands

Mnemonic: "NLU = Naturally Understands Language"

Question 4(c) OR [7 marks]

Explain stemming and lemmatization in detail. Also provide two examples of each.

Answer:

Definitions:

Process	Description	Method	Output
Stemming	Reduces words to root form by removing suffixes	Rule-based chopping	Word stem
Lemmatization	Reduces words to dictionary base form	Dictionary lookup	Valid word

Stemming Examples:

1. "running", "runs", "ran" → "run"
2. "fishing", "fished", "fisher" → "fish"

Lemmatization Examples:

1. "better" → "good" (comparative to base)
2. "children" → "child" (plural to singular)

Comparison Table:

Aspect	Stemming	Lemmatization
Speed	Faster	Slower
Accuracy	Lower	Higher
Output	May not be valid word	Always valid word

Mnemonic: "Stemming = Speed, Lemmatization = Language accuracy"

Question 5(a) [3 marks]

Define: 1) Word embeddings. 2) Machine Translation.

Answer:

Term	Definition	Purpose
Word Embeddings	Dense vector representations of words that capture semantic relationships	Convert text to numerical form for ML
Machine Translation	Automated conversion of text from one language to another	Enable cross-language communication

Key Features:

- **Word embeddings** preserve word relationships in vector space
- **Machine translation** maintains meaning across languages

Mnemonic: "Words become Vectors, Languages become Translations"

Question 5(b) [4 marks]

Explain Word2Vec in detail.

Answer:

Definition: Neural network technique that creates word embeddings by learning word associations from large text corpus.

Architecture Types:

Model	Description	Prediction
CBOW (Continuous Bag of Words)	Predicts target word from context	Context → Target
Skip-gram	Predicts context words from target	Target → Context

Process:

- 1. **Training:** Neural network learns word relationships
- 2. **Vector Creation:** Each word gets unique vector representation
- 3. **Similarity:** Similar words have similar vectors

Example: $\text{vector}(\text{"king"}) - \text{vector}(\text{"man"}) + \text{vector}(\text{"woman"}) \approx \text{vector}(\text{"queen"})$

Mnemonic: "Word2Vec = Words to Vectors through Context"


Question 5(c) [7 marks]

Consider that a manufacturer of a product has collected feedback from the customer and is now willing to apply sentiment analysis on it. What are the steps to be followed for the same? Explain in detail.

Answer:

Sentiment Analysis Pipeline:

Step	Description	Tools/Methods
Data Collection	Gather customer feedback	Surveys, reviews, social media
Data Preprocessing	Clean and prepare text	Remove noise, tokenization
Feature Extraction	Convert text to numerical	TF-IDF, Word embeddings
Model Training	Train sentiment classifier	Naive Bayes, SVM, Neural networks
Prediction	Classify sentiment	Positive/Negative/Neutral
Analysis	Interpret results	Generate insights and reports

Implementation Flow:**Preprocessing Steps:**

- **Remove special characters** and URLs
- **Convert to lowercase** for consistency
- **Remove stop words** (the, and, or)
- **Handle negations** (not good → negative sentiment)

Model Evaluation: Use metrics like accuracy, precision, recall, F1-score

Business Value: Understand customer satisfaction, improve products, identify issues

Mnemonic: "Collect, Clean, Extract, Train, Predict, Analyze = Ccetpa"

Question 5(a) OR [3 marks]

List out advantages of GloVe with respect to NLP.

Answer:

Advantage	Description
Global Context	Considers entire corpus statistics, not just local context
Linear Relationships	Captures semantic relationships through vector arithmetic
Efficiency	Faster training compared to Word2Vec on large datasets
Stability	Consistent results across multiple training runs

Key Benefits: Better performance on word analogy tasks, captures both local and global statistics

Mnemonic: "GloVe = Global Vector Excellence"

Question 5(b) OR [4 marks]

Explain challenges with TFIDF and BoW.

Answer:

Challenges Table:

Method	Challenges	Impact
TF-IDF	1. Ignores word order 2. Sparse vectors 3. No semantic similarity	Limited context understanding
BoW	1. Loses sequence information 2. High dimensionality 3. No word relationships	Poor semantic representation

Common Issues:

- **Vocabulary size:** Creates very large, sparse matrices
- **Out-of-vocabulary:** Cannot handle new words
- **Semantic gap:** "Good" and "excellent" treated as different

Solutions: Use word embeddings (Word2Vec, GloVe) for dense, semantic representations

Mnemonic: "TF-IDF and BoW = Sparse, No order, No semantics"

Question 5(c) OR [7 marks]

Consider that an e-mail service provider is willing to apply a SPAM detection technique. What are the steps to be followed to detect a SPAM e-mail? Explain in detail.

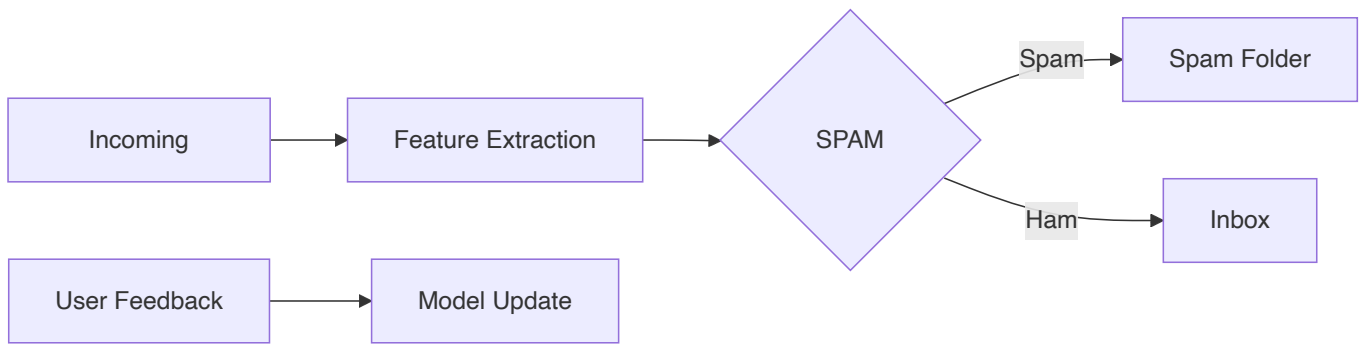
Answer:

SPAM Detection Pipeline:

Step	Description	Techniques
Data Collection	Gather labeled spam/ham emails	Email datasets, user reports
Feature Engineering	Extract relevant features	Subject analysis, sender patterns
Text Preprocessing	Clean email content	Remove HTML, normalize text
Feature Extraction	Convert to numerical form	TF-IDF, N-grams, metadata
Model Training	Train classifier	Naive Bayes, SVM, Random Forest
Validation	Test model performance	Cross-validation, test set
Deployment	Integrate with email system	Real-time classification

Feature Types:

Feature Category	Examples	Purpose
Content-based	Keywords, phrases, HTML tags	Analyze email body
Header-based	Sender, subject, timestamps	Check metadata
Behavioral	Sending patterns, frequency	Identify suspicious behavior

Implementation Diagram:**Model Evaluation Metrics:**

- **Precision:** Avoid false positives (legitimate emails marked as spam)
- **Recall:** Catch actual spam emails
- **F1-Score:** Balance between precision and recall

Continuous Learning: Update model with new spam patterns and user feedback

Mnemonic: "Collect, Engineer, Process, Extract, Train, Validate, Deploy = CEPTVD"